

Human Activity Recognition and Classifier Comparison for Support Vector Machines and 1-Dimensional Convolutional Neural Networks

Jake Sant¹ [117699M], Aiden Williams² [372001L], Ethan Zammit³ [4802L]

Department of Artificial Intelligence

University of Malta

*jake.sant.18@um.edu.mt*¹, *aiden.williams.19@um.edu.mt*², *ethan.zammit.19@um.edu.mt*³

Keywords—Human activity, Classification, Support Vector Machine, Convolutional Neural Network, Accuracy

I. INTRODUCTION

In this project we reviewed a number of published papers about human activity recognition and implemented a number of classical machine learning and deep learning classification techniques for the UCI dataset. Then, we collected our own data by performing different activities across multiple sessions. We then implemented the classification methods that worked best for the UCI dataset [1] on this new dataset. Finally, we compared the results for the different classifiers.

II. AIMS AND OBJECTIVES

The objective of this project was to implement a mix of deep learning and classical machine learning techniques to correctly classify the UCI dataset [1] by following the Anguita et al.'s methodology. Subsequently, a new dataset would be constructed which includes a number of additional activities not found in the UCI dataset. Using this new dataset, the previously implemented techniques would be adapted to classify these new activities while maintaining high performance.

III. BACKGROUND RESEARCH

A. Sensing Devices

When considering applications in Human Activity Recognition (HAR), a range of sensors may be applied to recognise the activities being performed. One option is to use external sensors such as a camera, a microphone or as used in [A system for change detection and human recognition in voxel space using the Microsoft Kinect sensor.] using the Microsoft Kinect as a 3D sensing device applied to Human Activity Recognition. However, these are usually limited to a static place, as the equipment to externally scan the activities needs a setup. On the other hand, there are wearable/portable sensing devices, such as accelerometers, gyroscopes or possibly heart

rate monitors, which also have their own challenges. There were even methods such as [Emmanuel Tapia, Stephen Intille, Louis Lopez, and Kent Larson. The design of a portable kit of wireless sensors for naturalistic data collection. *Pervasive Computing*, pages 117- 134] where the two types of sensing devices were combined, aiming to analyze intricate activities.

B. Data Acquisition

During the process of data collection, care should be taken to keep conditions as natural as possible, as large deviation from natural conditions might make the model unusable, as conditions may be drastically different. The variety must also be maintained as to aim for a robust model which can handle different conditions, noise levels, users and environments [Human Activity and Motion Disorder Recognition: Towards Smarter Interactive Cognitive Environments Jorge L. Reyes-Ortiz^{1,2}, Alessandro Ghio¹, Davide Anguita¹, Xavier Parra², Joan Cabestany², Andreu Catal'a²].

C. Signal Processing & Feature Extraction

A robust signal processing pipeline is essential to reliably convert raw inertial data into a format usable by a classifier, as the features from the data need to be carefully extracted and highlighted to allow the classifiers to distinguish between the activities using the reproducible characteristics, such as statistical results on windows of data. Data windowing is a very common practice when considering sensory data, as several statistical tools are only unlocked on bundled data, and allow for richer and denser information per entry. An element of overlap between windows is also often employed, as the overlap allows for smoother transitions between activities. [Human Activity and Motion Disorder Recognition: Towards Smarter Interactive Cognitive Environments Jorge L. Reyes-Ortiz^{1,2}, Alessandro Ghio¹, Davide Anguita¹, Xavier Parra², Joan Cabestany², Andreu Catal'a²]. Albeit the importance of reliable classification one also needs to consider the processing

time required for the signal processing pipeline, as in some applications real-time classifications are necessary, and if the feature vector is excessively large, performance may be heavily impacted.

D. Support Vector Machines

Support Vector Machines (SVM) are a type of classifier used for classification and regression. SVM seeks to classify samples as different classes on a hyperplane in multidimensional space. The classification process generates multiple hyperplanes to find the one which minimises an error value. Generally, support vector machines are used for binary classification problems. However, they can be adapted to multiclass problems by converting or reducing them into a set of multiple binary classification problems. However some samples cannot be classified linearly, hence SVMs can also make use of kernels to transform non-linear space into linear space, which was proposed by Boser et al. [2].

E. Convolutional Neural Networks

The Convolutional Neural Network (CNN) is a type of deep learning neural network model commonly applied for computer vision applications. Like other deep learning neural networks, CNNs consist of an input layer, any number of hidden layers, and an output layer. Unique to the CNN is the use of the convolutional layer, these layers use a kernel that goes over the input tensor, causing the convolution. Fully connected linear layers can achieve the same results however, for large inputs FC layers will have many weights. With convolutional layers the number of these weights is reduced. Each layer has its own activation function, which acts as a filter for whether a neuron should fire or not. In this project the RELU, Softmax and LogSoftmax activations were used. After a number of convolutional layers an FC layer is usually inserted, and in this project each CNN has an FC layer at the end.

IV. LITERATURE REVIEW

Anguit et al. [3] classified activities into two different categories: static and dynamic activities. Using a One-Vs-All approach and a Laplacian kernel, they obtained highly accurate rates of classification. Demrozi et al. [4] compiled a number of studies on human activity recognition (Human Activity Recognition using Inertial, Physiological and Environmental Sensors: A Comprehensive Survey) and found that out of 149 papers published between January 2015 and September 2019, 53 examined deep learning models and the remaining 96 were classical machine learning models. Cho et al. [5] propose a divide and conquer approach using a 1-dimensional convolutional neural network on the UCI dataset.

V. DESIGN AND METHODOLOGY

A. Data Collection

For the data collection, a barebones Android app was developed, which reads the triaxial accelerometer and gyroscope data at a constant rate of 50hz. Additionally, timestamps and a unique device identifier were appended to each entry, along with an option to label recording sessions, all to simplify the processing pipeline and avoiding the need to manually label all the data afterwards.

Using this app, several activities were carried out and recorded, whilst trying to vary the conditions as much as possible, such as putting it in the left and right pockets and having it upside down or upright, aiming to diversify the data.




Fig. 1: Our Data Distribution

Figure 1 depicts the different amounts of data collected per activity. Unfortunately, there was a considerable difference in the amounts of data, where some activities such as jumping had much less data than other activities like walking. This can lead to the accuracy of classification of data to be misleading, since if for example the classifier is completely random, and it guesses 'walking', it is more likely to be right than if it randomly guesses 'jumping', and since there is more walking data, it may seem more accurate than it really is.

B. Signal Processing

The JSON data exported from the mobile application was then ready to be imported by a python preprocessor. As is, the data is not usable by the classifiers, not only does it contain a lot of noise, but a single entry gives very little information about what is happening.

Consequently, several steps were taken to maximise the usability and the potential for classification of the data, such as synchronisation of gyroscope and accelerometer signals, use of a sliding window to group data, signal filtering and feature extraction using several statistical measures.

C. Synchronisation

Due to a limitation from the application, the gyroscope and accelerometer collections worked asynchronously, and thus

the entries were likely not to be aligned properly (a few milliseconds of difference). And if the phone screen was turned off, data collection would stop and restart when the screen is turned on, which would cause a substantial gap within the data of a session.

A moving average of the differences was calculated over all the data. If the average difference is below 10ms, the values would be considered as forming part of the same timestamp, whilst if the average value was larger, an entry was removed from the accelerometer and gyroscope data, until the average restabilized. This eliminates most outliers, and if there are antired out of synch by a substantial margin, they are removed. This was considered quite harshly, as several techniques such as Fast Fourier transform rely on constant frequencies to provide desirable outputs. The first and last 2 seconds were also removed for each session, as time needs to be accounted for whilst the user puts the device in his/her pocket, which would provide inaccurate data.

D. Sliding Window

As carried out in [<https://www.sciencedirect.com/science/article>] a sliding window of size 128 with a stride of 64 was adopted, which in simpler terms means grouping 128 entries (2.56 seconds of data) for each window, and overlap each window by 1.28 seconds.

This is one of the most crucial steps in data processing, as one single entry does not describe a whole lot about what a user is doing. Before any preprocessing, each entry simply has six values (two triaxial vectors), which are at a single point in time. When the sliding window is introduced, we group a window's worth of entries into one, out of which we get a representation of what happens over a period (2.56 seconds in our case), where each window will represent 128 raw entries, containing 6 features (triaxial data) each. Which when feature mapped, summarizes features of this two-dimensional data frame, back into a one-dimensional entry, whilst still retaining the fact that an entry has now the added dimension of time.

E. Signal Filtering

A set of transformations were applied to each window of entries, which were used to filter and prepare the data for the feature extraction stage.

1) *20hz Butterworth Filter*: A low-pass Butterworth filter with a corner frequency of 20hz, which removes most of the noise, as most changes in the acceleration happened at much lower frequencies, and through this separation, the most relevant changes in acceleration were kept. In fact, according to [D.M. Karantonis, M.R. Narayanan, M. Mathie, N.H. Lovell,

and B.G. Celler. Implementation of a real-time human movement classifier using a triaxial accelerometer for ambulatory monitoring. IEEE Transactions on Information Technology in Biomedicine, 10(1):156–167, 2006.] most bodily movements are contained below 15Hz

2) *0.3hz Butterworth Filter*: Another low pass Butterworth filter was applied with a corner frequency of 0.3hz, which filter separates gravitational acceleration from body acceleration, since changes in gravitational acceleration happen relatively slowly, a low corner frequency can distinguish and extract these changes as gravitational acceleration.




Fig. 2: The Acceleration in a single window

Figure 2 shows the different forces of acceleration experienced whilst the user was walking. The Body forces show the filtered acceleration when separated from the gravitational acceleration. At idle, these forces should be zero. The strides of the user's walking may also be observed as fluctuations in the green line, quickly dropping when the wearer's foot descends and hits the floor. The other three stable plots are the gravitational forces experienced, these should not change as much, as the gravitational force experienced does not fluctuate quickly. At idle the vector of these three forces should point downwards with a magnitude of $\sim 9.81m/s^2$.

3) *The Jerk*: Jerk is defined as the change in acceleration and was calculated by differentiating the acceleration values by time: $\delta t = 1/50hz = 0.02s$. This gives a better idea of how drastically the accelerations change, such as when one is jogging and the leg where the phone is sitting stops sharply as it hits the ground.

4) *The Magnitude*: The magnitude is the measure of how much the values differ from zero. This comes useful when one simply needs an idea of the amount of force applied, without considering the effects of directionality.

5) *Fast Fourier Transform*: FFT enabled the calculation of frequency components based on the time-varying signals.

When considering Human Activity, one can note that there is a lot of repetitive patterns, in which case transformations such as FFT are used to calculate the discrete Fourier transform, and give a rich statistic for summarizing a window of these repetitive signals.

<https://www-proquest-com.ejournals.um.edu.mt/docview/230175170/fulltext/1965/4312904/409834859/1?accountid=230175170>

F. Feature Extraction

Excluding the label for each activity, each session in the dataset contains 589 features. These features were extracted by applying a number of different statistical measures to the different extracted signals in each sliding window.

Each signal has the statistical measures in Table 1 applied to them.

Statistical Measure	Description
mean	Mean value in the window
min	Smallest value in the window
max	Largest value in the window
std	Standard deviation
entropy	Maximum Entropy Distribution
mad	Median Absolute Deviation
iqr	Inter-quartile Range
energy	Sum of each squared value in the window
sma	Signal magnitude area, sum of absolute values in the window
arCoeff	Coefficients of the autoregressive model using Burg's algorithm
correlation	Correlation between X, Y and Z signals
angle	Angle between 2 vectors of signals

Table 1

The statistical measures in Table 2 are applied only to FFT signals.

Statistical Measure	Description
band energy	Energy value bounded by bins of the set of FFT frequencies
maxInds	Fast Fourier Transform bin which matched the highest value in the window
skewness	Fisher-Pearson coefficient of skewness
kurtosis	Kurtosis of the array of values in the window
meanFreq	Weighted average of the signals in the window

Table 2

G. Signal Processing

H. Support Vector Machine Classifier

1) *Data Pre-Processing*: Firstly, the training and testing datasets were converted to Pandas dataframes and the labels for each activity were separated into their own variables. Using a LabelEncoder, each activity label is converted into a numerical value.

2) *Comparing Different Classification models*: In total, four different classic machine learning classifiers were used on the dataset developed by Anguita et al [3]. These classifiers are Gaussian Naïve Bayes, AdaBoost, Stochastic Gradient

Descent and a Support Vector Machine. These were trained and tested using the respective datasets, and their accuracy, F-beta, precision and recall scores were recorded.

Classifier	Accuracy	F-Beta	Precision	Recall
Gaussian Naïve Bayes	0.7134	0.7252	0.7555	0.7134
AdaBoost	0.4065	0.2520	0.4289	0.4065
Stochastic Gradient Descent	0.9600	0.9603	0.9605	0.9600
Support Vector Machine	0.9655	0.9657	0.9659	0.9655

Table 3: Performance scores of each classifier

As can be seen in Table 3, the Stochastic Gradient Descent classifier and the Support Vector Machine had the highest scores, with each result being over 0.96. However, the SVM's performance metrics had slightly higher results. Furthermore, the cost function in Stochastic Gradient Descent is based on a single set of samples in the training data and its positive results do not translate to accurate results in the testing phase. It is for this reason that the Support Vector Machine classifier was chosen to classify the UCI dataset [1] and, later on, the dataset built by ourselves.

I. Classification

The RBF kernel is defined as the exponential function $\exp(-\gamma|x - x'|)^2$ [6], where x and x' are two feature vectors, and γ is the gamma parameter in the classifier. Gamma's value is scale, meaning that the parameter is the reciprocal of the number of features multiplied with the variance of the input data. For this implementation, we opted for a One-Vs-All approach. A One-Vs-All approach divides the data points into just two classes: a certain activity and the other classes. Therefore records labelled as *SITTING* are a single class, and the other activities are treated as having a single label.

J. Convolutional Neural Networks

From these, Cho et al. [5] presented the choice and implementation of a CNN use for HAR the best. Hence, the implementation described in this report is based on the architectures described. Unlike Cho et al. we did not implement a first stage dynamic-static split model and instead opted to split the labels manually. The Pytorch [reference it] library was used to implement the CNNs, and the sklearn [reference it] library to evaluate the results.

1) *Data Split*: The train/valid/test split was an 80/20 train/test split, then 80/20 train/validation split for the UCI dataset, and a 90/10 train/test split, then 90/10 train/validation split for Our Dataset.

2) *Dataset Object*: To facilitate this dynamic/static split we defined a Python dictionary with each label having a number assigned to it, starting from 0. The numbering needs to start from 0 as the output of neural networks implemented

in Pytorch always start from 0. Another requirement for implementing neural networks in Pytorch is to implement a custom Dataset object. The function of this object is to read the data from a source and define its X, data and Y, labels counterparts. The Pandas [reference it] library was used at this stage due to its use of the numpy [reference it] library and efficient data management. It is important that the X component is in the shape: `length(data.columns), 1, length(data.rows)`.

The dataset object was initialised 3 times, for the Train, Validation and Testing data.

The pie charts included below show the training, validation and testing label distribution in order.




Fig. 3: UCI Dynamic Label Distribution




Fig. 4: UCI Static Label Distribution



Fig. 5: OD Dynamic Label Distribution




Fig. 6: OD Static Label Distribution

In development, it was noted that the data distribution of the dynamic activities in our dataset was unequal. This was causing issues with the CNN's accuracy. To combat this added functionality was added with the aim of increase the quality of the model's output. In this project the number of rows for each eligible label was stored. Using the functionality provided by the Pandas library, the index for each label was stored in a list

linked to its respective label. Using Python list slicing these lists were cut down to their lowest common length. These indices were then converted into a new Pandas dataframe and this balanced dataset was used for the dynamic model.




Fig. 7: OD Dynamic Training Balanced Label Distribution

Finally, the dataset objects are used to initialise a Dataloader object. This object gives us the functionality of delivering batches as inputs to the models at training and testing. It was found that a batch size of 32 worked best for the Static models, and a batch size of 64 worked best for the dynamic models.

3) *Model Creation:* The CNN models were implemented following the design of the below diagrams. Every model used the Cross-Entropy as their loss function and the ADAM optimizer. The learning rate for the UCI Dataset Models and Our Dataset Static model was set at 0.0005 and were trained for 10 epochs. The learning rate for the Our Dataset Dynamic model was set at 0.00005 and were trained for 11 epochs.

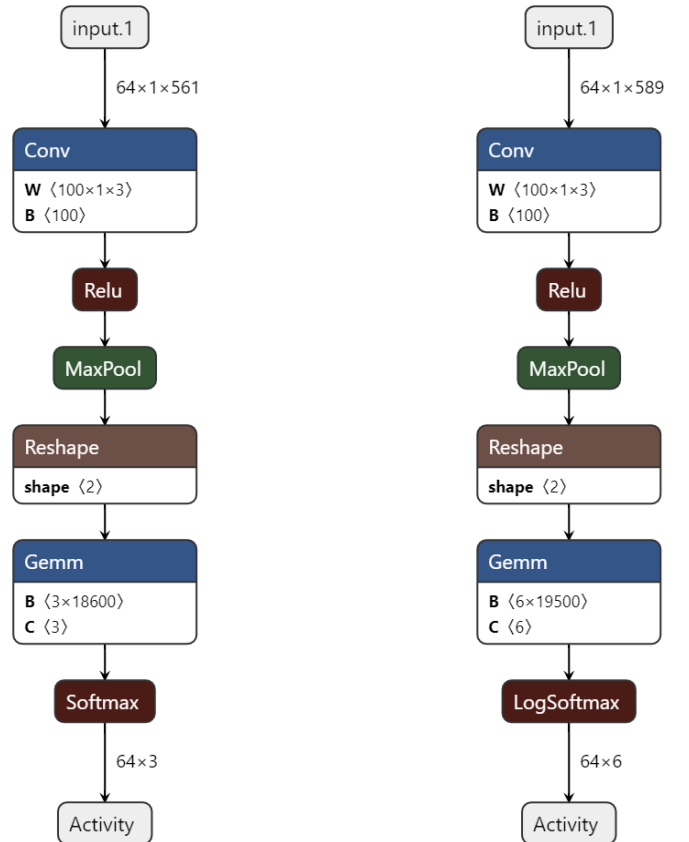


Fig. 8: UCI(left) and OD(right) Dynamic Model Structure

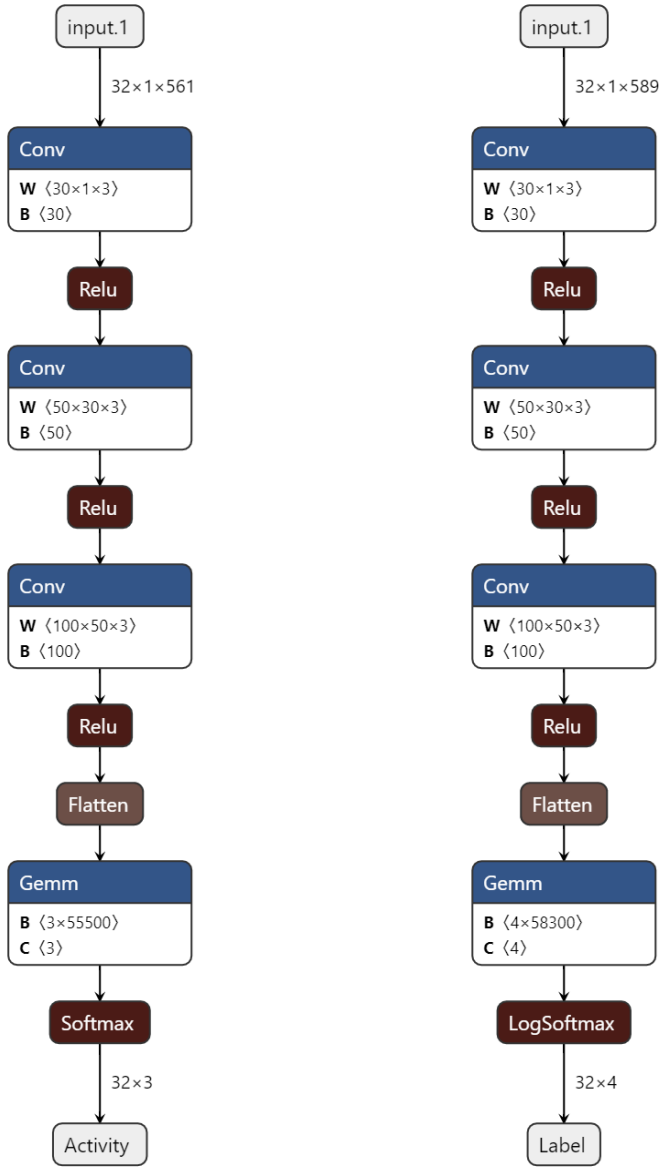


Fig. 9: UCI(left) and OD(right) Static Model Structure

As seen in figures 9 and 8 the four models follow the a similar structure having a Convolutional layer that is then followed by a fully connected layer. The OD models differentiate at the final activation function where it was found that the Softmax activation function was not activating properly. Seeing this another activation function was used, the LogSoftmax function. It is also worthwhile to note that the final layer fully connected layer is different since the UCI dataset and Our dataset have a different number of total features. These models accept a tensor of the shape: (batchsize, 1, number_of_features). The output of item in this batch is a number that corresponds to the encoded label. Dropout is present at the fully connected layer, set to 50% to combat overfitting. Dropout is useful since it applies a 50% chance of ignoring a neurons outputting, which

forces the model to learn redundancy.

VI. EVALUATION

A. Evaluation Metrics

In order to determine which classifier provided the most accurate predictions, four different performance metrics were recorded and acted as a score. These metrics are **accuracy**, **f-beta**, **precision** and **recall**.

Accuracy is the ratio of the true labels y on the set of predicted labels y' , where

$$accuracy(y, y') = \frac{1}{n_{samples}} \sum_{i=0}^{n_{samples}-1} 1(y'_i = y_i)$$

Precision is the ratio of true positives during prediction, and it is calculated as

$$P = \frac{T_P}{T_P + F_P}$$

where T_P is the number of true positives and F_P is the number of false positives.

Recall is the ratio of correctly identified positive labels, and it is calculated as

$$R = \frac{T_P}{T_P + F_n}$$

where T_P is the number of true positives and F_n is the number of false negatives.

Beta (or **F-Beta**) is calculated based on the precision and recall scores of the classifier, wherein precision is multiplied by some parameter γ , thereby giving more importance to the precision value. This is calculated as

$$F_\gamma = (1 + \gamma^2) * \frac{P * R}{(\gamma^2 * P) + R}$$

For evaluation purposes, a γ value of 0.5 was used.

B. Evaluating the SVM

Figure X is a confusion matrix representing the accuracy of the classification of each activity as a percentage. The leading diagonal represents correctly labelled activities. Each activity group classified had an accuracy of 96% or higher, with half activities being classified at an accuracy of 99-100%. 1% of driving activities were misclassified as sitting. This was to be expected as during driving sessions there were times (i.e. being stuck in traffic) where the driver was stationary and in the same position as that of someone seated down. 2% of laying activities were misclassified as sitting, this is understandable as the two stationary activities have similar positions.

C. Evaluating the CNNs

The CNN models were evaluated during training as well as after testing. During training, the training and validation




Fig. 10: Fig. X: Confusion Matrix for Support Vector Machine

loss was printed. When a hyperparameter was being amended during the development the shift in the loss output would indicate the success or failure of the amendment. In general, when the validation loss is more than the training loss, the model is considered to be overfitting, when the validation loss is less than the training loss, the model is considered to be underfitting. It was our aim to set the validation and training losses as equal to each other as possible. The hyperparameters set, as described in the methodology are a result of this process. After the testing steps were completed, the same evaluation metrics used for the classical ML techniques were implemented.

1) *UCI Dynamic & Static*: These models achieved very high results (more than 90%) however a little overfitting was observed. To combat this overfitting, the epoch count was reduced however this started affecting general accuracy and so it was decided to keep epochs at 10 and accept a little overfitting. It was also noted that although results were mostly consistent, each new run would change the overall accuracy score by 2% on average. The Static model achieved the highest results from the CNN models. Due to their high accuracy, the f-score, precision and recall metrics resulted in similarly high results and offered little new information.

2) *Our Dataset Dynamic & Static*: Initially the architecture for the UCI Dataset models would not function on our dataset. It was found that the activation function of the last fully connected layer had to be changed from Softmax to LogSoftmax. This affected the performance of the models as it increased the processing load. The Dynamic model achieved good results with an accuracy of more than 60%, when considering it




Fig. 11: Dynamic(left) and Static(right) UCI Model Confusion Matrices

trained on a reduced sized dataset with an equal distribution of 15% for each label. The extra effort put in to balance out the training distribution was expected to benefit the model in terms of accuracy, however there was no difference between the results of the model before the data balancing and after. The higher precision score of 71% compared to the recall of 61% shows that the model is still able to predict the relevant label. The Static model achieved very high results (more than 90%). The same steps taken for the UCI dynamic & static models were taken and decided upon.

VII. CONCLUSION AND FUTURE WORK

In this project it was found that the SVM classifier achieved the best results from all the techniques used for both the UCI dataset and our dataset when considering all activities. As shown in the evaluation section, the CNN implemented for static activity classification achieved near equal results. Considering the fact that the CNN implementations required more time and effort to actually function as HAR classifiers, the results obtained are satisfactory. Based on the results

obtained, and the quality of the data constructed, the SVM classifier has shown to be the better performing classifier and ideal for HAR.

VIII. DISTRIBUTION OF WORK

For this project each of us chose a section to be the main maintainer of. From there we distributed tasks and worked on the entire project together. Below we have listed the section owners.

- Jake Sant - Support Vector Machine & Feature Extraction
- Aiden Williams - Convolutional Neural Network & Classical Machine Learning Comparison
- Ethan Zammit - App Development & Data Processing

REFERENCES

- [1] D. Anguita, A. Ghio, L. Oneto, X. Parra, and J. L. Reyes-Ortiz, "A Public Domain Dataset for Human Activity Recognition Using Smartphones," in *21st European Symposium on Artificial Neural Networks, Computational Intelligence and Machine Learning*, Bruges, 2013. [Online]. Available: <https://www.elen.ucl.ac.be/Proceedings/esann/esannpdf/es2013-84.pdf>
- [2] B. E. Boser, I. M. Guyon, and V. N. Vapnik, "A training algorithm for optimal margin classifiers," in *Proceedings of the fifth annual workshop on Computational learning theory - COLT '92*. New York, New York, USA: ACM Press, 1992, pp. 144–152. [Online]. Available: <http://portal.acm.org/citation.cfm?doid=130385.130401>
- [3] D. Anguita, A. Ghio, L. Oneto, X. Parra, and J. L. Reyes-Ortiz, "Human activity recognition on smartphones using a multiclass hardware-friendly support vector machine," *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*, vol. 7657 LNCS, pp. 216–223, 2012. [Online]. Available: http://link.springer.com/10.1007/978-3-642-35395-6_30
- [4] F. Demrozi, G. Pravadelli, A. Bihorac, and P. Rashidi, "Human Activity Recognition Using Inertial, Physiological and Environmental Sensors: A Comprehensive Survey," *IEEE Access*, vol. 8, pp. 210 816–210 836, 2020.
- [5] H. Cho and S. M. Yoon, "Divide and Conquer-Based 1D CNN Human Activity Recognition Using Test Data Sharpening," *Sensors*, vol. 18, no. 4, 2018. [Online]. Available: <https://www.mdpi.com/1424-8220/18/4/1055>
- [6] B. Schölkopf, K. Tsuda, and J.-P. Vert, "A primer on kernel methods," in *Kernel Methods in Computational Biology*, B. Schölkopf, K. Tsuda, and J.-P. Vert, Eds. The MIT Press, jul 2004, pp. 35–70. [Online]. Available: <https://doi.org/10.7551/mitpress/4057.001.0001>